

Stage III: Predicting NYPD Misconduct Case Outcomes and Penalties

Helen Wang

Literature Review

In 2020 alone, 53.8 million U.S. residents reported having contact with law enforcement (Tapp & Davis, 2020). While the majority (88%) of respondents reported having satisfactory interactions with police, about 2% or 1 million respondents reported having negative experiences, characterized by the threat and/or use of force as well as other forms of misconduct (Tapp & Davis, 2020). Police misconduct, defined by the New York Police Department (NYPD) Civilian Complaints Review Board, is any action or charge that may result in an officer being subjected to an administrative trial process (Cubitt et al., 2022). Conduct that falls under this category includes, but is not limited to, improper use of force, abuse of authority, discourteous behavior, and offensive language (Cubitt et al., 2022). In recent years, police misconduct has become a prominent issue in the eyes of the American public, especially given racial/ethnic disparities in experiences, with numerous high-profile cases resulting in the death of many lives. Releases of police misconduct records in recent years have allowed researchers to better characterize and understand the escalating problem at hand.

Research has shown that there are several correlates of police misconduct. In general, male officers are more likely to get complaints than their female counterparts (Harris & Worden, 2014). Younger, less experienced officers are also at higher risk along with officers with military experience (Harris & Worden, 2014). Officers who are more productive (i.e. with more arrests and stops) and who have more civilian interaction are also more likely to accrue greater complaints of police misconduct (Harris & Worden, 2014; Rozema & Schanzenbach, 2019). Some reports have also suggested that officers from minority populations may also be more likely to commit misconduct, though other studies have suggested that this may be a result of differential task or geographic assignment to high-crime precincts (Cubitt et al., 2022; Harris & Worden, 2014)

Police misconduct records only represent a small portion of actual misconduct cases; it is estimated that approximately only a third of complaints end up being actually filed (Harris & Worden, 2014). Of those filed, legal and institutional barriers make it difficult for complaints to result in disciplinary measures. Reports have shown that only an eighth of civilian-initiated complaints are sustained, with most complaints often being declared as either exonerated (act verified but found to be proper), unfounded, or not sustained (insufficient evidence) (Harris & Worden, 2014). If sustained, about 1 of 24 cases result in

sanctions for officers involved, with most sanctions often not commensurate with the misconduct that occurred (Harris & Worden, 2014).

Past work looking into police misconduct at various police departments across the nation has identified significant patterns. Rozema & Schanzenbach (2019), while examining police misconduct in Chicago, found that officers with moderate numbers of misconduct allegations were at no greater risk of committing serious misconduct than officers who had no misconduct allegations at all. Instead, they found that officers who were in the top 1% quartile of misconduct allegations were more likely to commit serious misconduct, generating almost 5 times the number of payouts and 4 times the total damage payments. When complaints against officers are sustained, almost 36% of offending officers accrue another sustained complaint at some point during their career (Harris & Worden, 2014). The risk of obtaining another complaint was found to be higher in the first months succeeding the first complaint but dropped significantly afterward. Notably, officers who are sanctioned for the complaint are not only more likely to engage in misconduct but do so more rapidly (Harris & Worden, 2014). Overall, findings suggest that a small subset of repeat offenders are responsible for a large portion of police misconduct reported, encompassing over \$1.5 billion in lawsuit settlements across the nation (Cubitt et al., 2022).

Cubitt et al. (2022) also noted that differences in case outcomes depend on officer and complainant characteristics. Female officers, for example, while less likely to accrue misconduct allegations in general than their male counterparts, were more likely to be sanctioned with remedial management action. Black and Hispanic civilians who submitted a complaint of police misconduct were 4.7 and 1.6 more likely to receive a not sustained ruling compared to White citizens (Headley et al., 2020). One study also found that racial mismatches between officer and complainant were linked to differing case outcomes (Wright II, 2020). The study found that Black complainants were more likely to receive a sustained ruling when misconduct was alleged against a white officer. On the other hand, white complainants were less likely to receive a sustained ruling when alleging misconduct against a black officer. These results differed across city departments, however, suggesting geographic differences.

Research Question(s)

The issue of police misconduct is important, especially given rising cases of police violence, police brutality, and fatalities as a result of such misconduct. Such issues have resulted in the loss of multiple lives, increased racial tensions across the US, and fractured public trust in law enforcement. The literature has explored various factors linked to police

misconduct and has similarly evaluated predictors of not only future misconduct but also misconduct case outcomes.

Few studies, however, have examined the impact of media coverage on police misconduct. Those who do have primarily limited their analysis to the impact of media coverage on public perceptions of police misconduct (Chermak et al., 2006; Dowler & Zawilski, 2007). No article to date has assessed the role that the media can play in influencing misconduct case outcomes, despite the potential pressure that media can have on ensuring that sufficient sanctions are levied.

To address this gap, in this study, we will utilize a dataset of police misconduct cases drawn from the New York City Police Department (NYPD) to do the following:

1. Assess the impact that media coverage/visibility of a given police misconduct incident has on its case outcome, specifically its complaint disposition, as well as the type of resulting penalty.
2. Using the literature, identify and verify additional factors predictive of police misconduct case outcomes such as officer/victim race, gender and age, along with incident location, police contact type, etc.
3. Develop a supervised model to predict case and penalty outcomes for police misconduct cases.

Data

To answer these questions, I will utilize police misconduct cases from the NYPD. Established in 1845, NYPD is one of the oldest and largest police departments in the nation, encompassing over 36,000 officers and 19,000 civilian employees (*About NYPD*, n.d.). Across its 78 precincts, the department serves over 8.5 million different individuals. The Civilian Complaint Review Board (CCRB), which separated from NYPD in 2000, has compiled a database of over 395,000 police misconduct cases from 2000 to 2025 (*Civilian Complaint Review Board (CCRB) Database*, n.d.). The database contains 4 datasets:

1. Allegations Against Police Officers: a list of all closed allegations made against NYPD officers, including information about the complainant, the officer, allegation, and resulting deposition (*Civilian Complaint Review Board: Allegations Against Police Officers | NYC Open Data*, n.d.)
2. Complaints Against Police Officers: a list containing information such as dates, locations, and circumstances surrounding the allegation (*Civilian Complaint Review Board: Complaints Against Police Officers | NYC Open Data*, n.d.)

3. Police Officers: a list of all NYPD officers and the number of total and substantiated complaints on their record (*Civilian Complaint Review Board: Complaints Against Police Officers* | NYC Open Data, n.d.)
4. Penalties: a list containing case and trial penalty information (*Civilian Complaint Review Board: Penalties* | NYC Open Data, n.d.)

To get media coverage information, I will use a subset of data available from the Mapping Police Violence Project, which used google alerts to get news articles on police violence events to construct their dataset (*Mapping Police Violence*, n.d.-a). Currently, the Mapping Police Violence Project has records of 48 separate police violence incidents, with associated news article links, that occurred between 2013 and 2024 (*Mapping Police Violence*, n.d.-b). Given the sample size, I will also utilize the New York Times API to scrape for relevant news articles on the officers and their associated misconduct cases.

Preprocessing and Descriptives

Given the sheer number of datasets I am utilizing for this study, merging datasets effectively is a critical component. I first merged the Allegations Against Police Officers and Complaints Against Police Officers datasets together by inner join, as there are multiple allegations forming single complaints (i.e. allegation ID is clustered by complaint ID). I then merged the resulting dataset to Penalties by left join, since Penalties are only given to complaints that have been ruled as “Sustained” and I want to keep non-sustained case outcomes. Finally, I merged this result with Police Officers by inner join. Police Officers contain all officers on the NYPD roster, regardless if they have a complaint against them. For this study, as I am interested in exploring case outcomes and resulting penalties, I used inner join to drop all officers who have not been subject to a complaint/allegation. This left me with a resulting dataset of 235,939 instances.

As this dataset lacked any variables for media coverage, I utilized the Mapping Police Violence Project dataset. After filtering for instances from the NYPD and for cases where the officers involved were known, I merged by officer name and year of incident using left join. Because the Mapping Police Violence Project dataset runs from 2013-2024 and the NYPD CCRB runs from 2000 to 2025, I elected to filter for instances that occurred after 2013. I also filtered for cases that occurred before 2020 to avoid potential endogeneity from COVID as well as the Black Lives Matter Movement (which was sparked in part by several police brutality and murder incidents that occurred around the time, which would likely skew media coverage).

For our main target variables, I am focusing on case outcome (CCRB Complaint Disposition) and resulting penalty (NYPD Officer Penalty). I utilized the CCRB Complaint

Disposition rather than the individual allegation dispositions as (1) only Sustained Complaints rather than Sustained Allegations result in a penalty and (2) multiple allegations are clustered under individual complaints. I also used the NYPD Officer Penalty instead of other penalty variables as they are typically recommended penalties and NYPD has final authority on which penalties are actually implemented.

Upon examination of the CCRB Complaint Disposition variable, in addition to expected disposition types (Substantiated, Unsubstantiated, Exonerated, Unfounded) there were also miscellaneous types such as “Complainant Uncooperative”, “Complaint Withdrawn”, “Subject Resigned”, etc. For the miscellaneous types, as they effectively meant the disposition did not happen, I recoded them as NaN and dropped those rows, resulting in a dataset of 36,666.

The NYPD Officer Penalty variable was also highly variable and in string format. In terms of penalties, some officers received multiple types and for different durations. I dummy coded this variable to capture the types of penalties they received, which may result in a loss of granularity but keeping it in its original state would have been difficult for analysis.

Using only the Mapping Police Violence Project, there are only roughly 33 instances where there was media coverage over the period of interest. Due to the scarcity of instances that fell under this criteria, I also incorporated data from the NYT Article Search API. Given rate limits, I elected to take a random subset of 4000 rows from the merged dataframe and extracted unique officer names and years. Using the officer name and the search terms “AND (police OR officer OR NYPD) AND (misconduct OR force OR brutality OR violence)” while filtering for articles within a year of the misconduct instance, I scraped for any hits. I then inner joined the results to the merged dataframe, giving me a resulting dataset of 10322, with the numbers reflecting the fact that multiple allegations are being clustered under singular complaints. While smaller than the full dataset, scraping hits for all misconduct cases would have been impossible—as such, I elected to look only at 4000 randomly selected cases.

Furthermore, I also utilized z-score normalization for numerical data variables and (depending on data distribution), reset intervals for categorical variables prior to dummy coding. After cleaning, we had minimal missing data: Incident Hour (0.4%), Precident of Incident Occurrence (1.7%), and Incident Month (0.1%). Data was imputed using MICE.

With the reduced dataset, we examined the distributions for our target variables below:

Table 1. CCRB Complaint Disposition


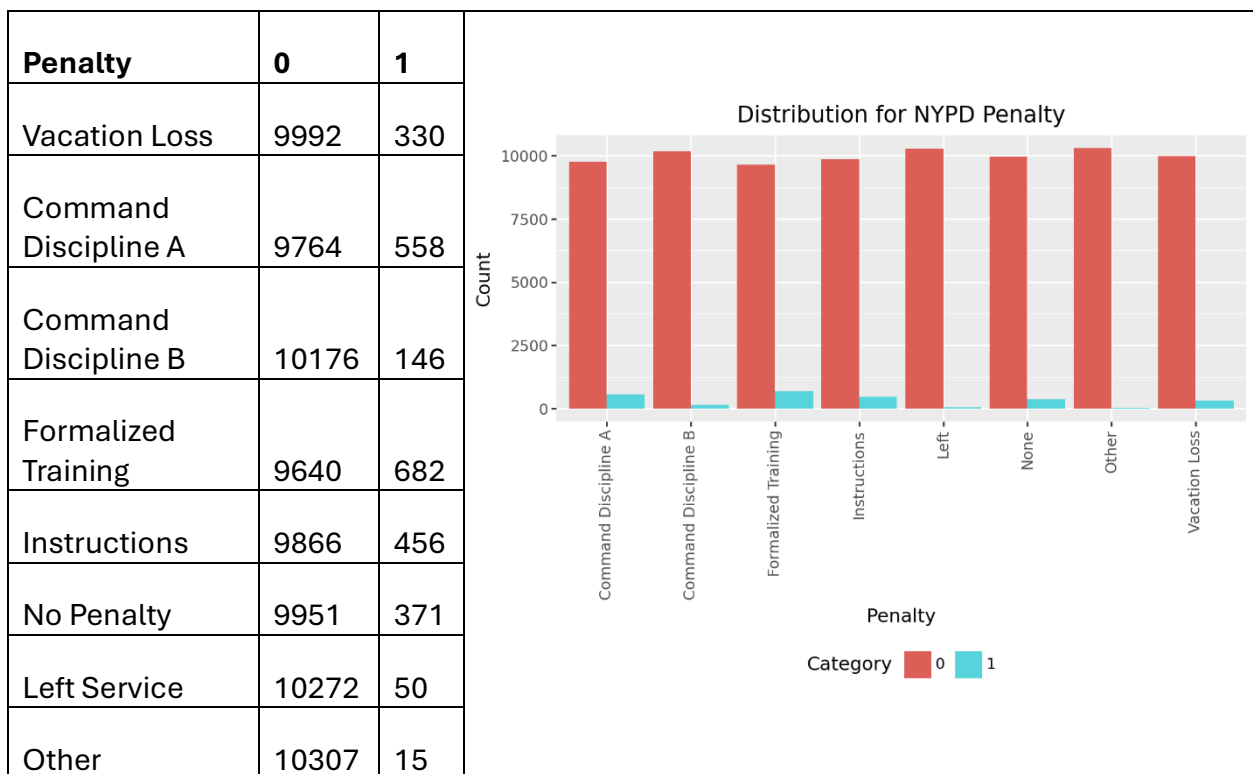


Table 2. NYPD Officer Penalty



Given the imbalance present in our target variables, I utilized synthetic minority oversampling technique (SMOTE) on training sets prior to running models..

Methodology

My main feature is a dummy variable indicating whether there was media coverage or not. Additional features that will be assessed include but are not limited to: Officer Days on Force at the Incident, Precinct of Incident Occurrence, Officer's Total Complaints, Officer Rank at Incident, Type of Misconduct, Victim demographics, Officer demographics, Location of Incident, etc.

Our main target variables in this case are the categorical variables: CCRB Complaint Disposition and NYPD Officer Penalty. For the first target variable, models need to be classifiers that are optimal for predicting multiclass variables and ideally are not sensitive to imbalance (given that is our biggest problem with our dataset). To account for imbalance in model performance, I will utilize F1-scores (weighted and macro) as well as balanced accuracy to assess model performance. Simultaneously, we will use confusion matrixes to assess model performance for each class.

For our second target variable, NYPD Officer Penalty, given that officers can receive multiple penalties (i.e. loss of vacation and formalized training), we need models that are optimal for multilabel classification. In this case, we assume that penalties are independent and elect to run separate binary classifiers for each label.

In both cases, I will avoid using models such as Decision Tree, which may be especially sensitive to imbalanced data. Simultaneously, I will also avoid models such as Naïve Bayes, given that some of the model assumptions (feature independence) may be violated given my dataset. Instead, I will utilize models such as logistic regression, support vector machine, k-Nearest Neighbors, Random Forest, and Gradient Boost. After splitting data into a 60-20-20 split for training, validation, and testing, I will run initial models. Depending on model performance, I will tune hyperparameters for a select subset of high-performing models on the validation set using randomized search and cross-validation.

Given that some of the models may perform better for smaller and low-dimensional datasets, I may have to reduce features using PCA. Other methodological considerations include dichotomizing my target variables (specifically for the CCRB Complaint Disposition).

Initial Results

We ran 5 different initial models to predict case outcome. Only two of these models (both ensemble method models) showed performance that was better than chance when utilized against the validation set. We report the F1-scores (weighted and macro), balanced accuracy, and confusion matrixes for each model below:

Table 3. Case Outcome Model Performance

Model	Metric	Score	Confusion Matrix
Logistic	F1 (Weighted)	0.298	
	F1 (Macro)	0.211	
	Balanced Accuracy	0.304	
SVM	F1 (Weighted)	0.296	
	F1 (Macro)	0.196	
	Balanced Accuracy	0.302	
KNN	F1 (Weighted)	0.526	
	F1 (Macro)	0.448	

	Balanced Accuracy	0.527	<div><p>Confusion Matrix</p><table><tr><th></th><th>0</th><th>1</th><th>2</th><th>3</th></tr><tr><th>0</th><td>82</td><td>27</td><td>22</td><td>28</td></tr><tr><th>1</th><td>109</td><td>330</td><td>50</td><td>106</td></tr><tr><th>2</th><td>17</td><td>11</td><td>58</td><td>13</td></tr><tr><th>3</th><td>155</td><td>176</td><td>107</td><td>361</td></tr></table></div>		0	1	2	3	0	82	27	22	28	1	109	330	50	106	2	17	11	58	13	3	155	176	107	361
	0	1	2	3																								
0	82	27	22	28																								
1	109	330	50	106																								
2	17	11	58	13																								
3	155	176	107	361																								
Random Forest	F1 (Weighted)	0.866	<div><p>Confusion Matrix</p><table><tr><th></th><th>0</th><th>1</th><th>2</th><th>3</th></tr><tr><th>0</th><td>71</td><td>12</td><td>2</td><td>74</td></tr><tr><th>1</th><td>3</td><td>537</td><td>2</td><td>53</td></tr><tr><th>2</th><td>0</td><td>0</td><td>66</td><td>33</td></tr><tr><th>3</th><td>7</td><td>19</td><td>4</td><td>769</td></tr></table></div>		0	1	2	3	0	71	12	2	74	1	3	537	2	53	2	0	0	66	33	3	7	19	4	769
		0		1	2	3																						
	0	71		12	2	74																						
1	3	537	2	53																								
2	0	0	66	33																								
3	7	19	4	769																								
F1 (Macro)	0.792																											
	Balanced Accuracy	0.744																										
Gradient Boost	F1 (Weighted)	0.870																										
	F1 (Macro)	0.810																										

	Balanced Accuracy	0.776	<div><div>Confusion Matrix</div><table><tr><th>True \ Predicted</th><th>0</th><th>1</th><th>2</th><th>3</th></tr><tr><th>0</th><td>80</td><td>6</td><td>4</td><td>69</td></tr><tr><th>1</th><td>3</td><td>527</td><td>2</td><td>63</td></tr><tr><th>2</th><td>0</td><td>0</td><td>75</td><td>24</td></tr><tr><th>3</th><td>19</td><td>13</td><td>4</td><td>763</td></tr></table></div>	True \ Predicted	0	1	2	3	0	80	6	4	69	1	3	527	2	63	2	0	0	75	24	3	19	13	4	763
True \ Predicted	0	1	2	3																								
0	80	6	4	69																								
1	3	527	2	63																								
2	0	0	75	24																								
3	19	13	4	763																								

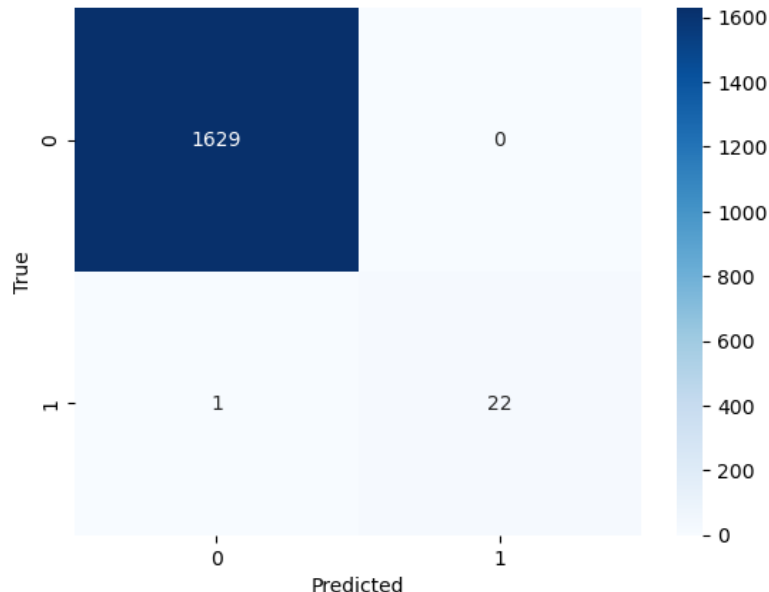
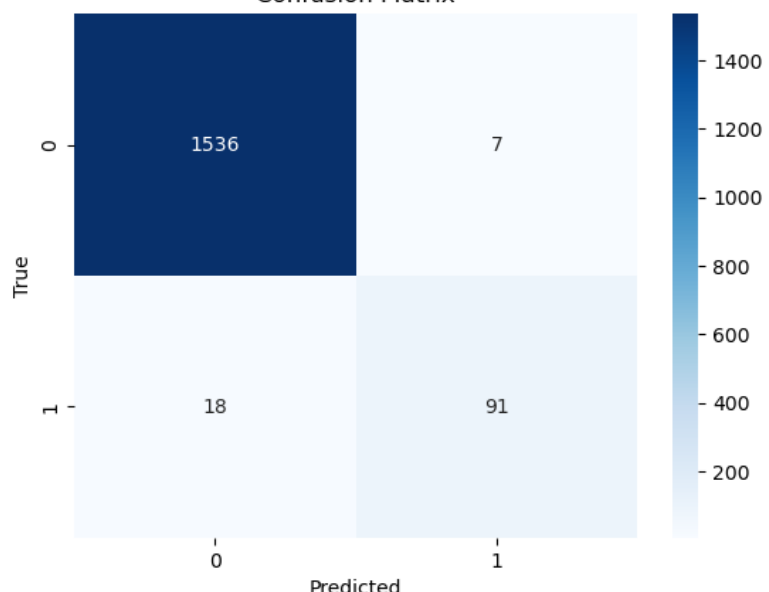
All non-ensemble method models appear to perform poorly, exhibiting low overall accuracy, precision, and recall. Examination of the confusion matrixes suggests that models are doing well in predicting "Unsubstantiated" (Class 3) cases but struggle with the other categories. Random Forest and XG Boost appear to be moderately successful with predicting "Substantiated" (Class 1) in addition to "Unsubstantiated" (Class 3) cases but also struggle with "Exonerated" and "Unfounded" cases (specifically for recall rather than precision). The two classes, "Exonerated" and "Unfounded," were the primary source of imbalance in our dataset, being the minority classes. Despite efforts to correct this via SMOTE, it appears that it still affected model performance.

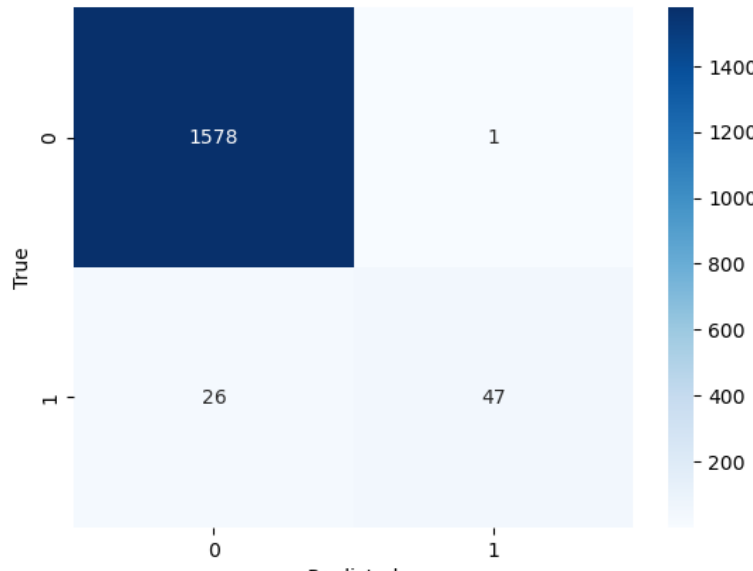
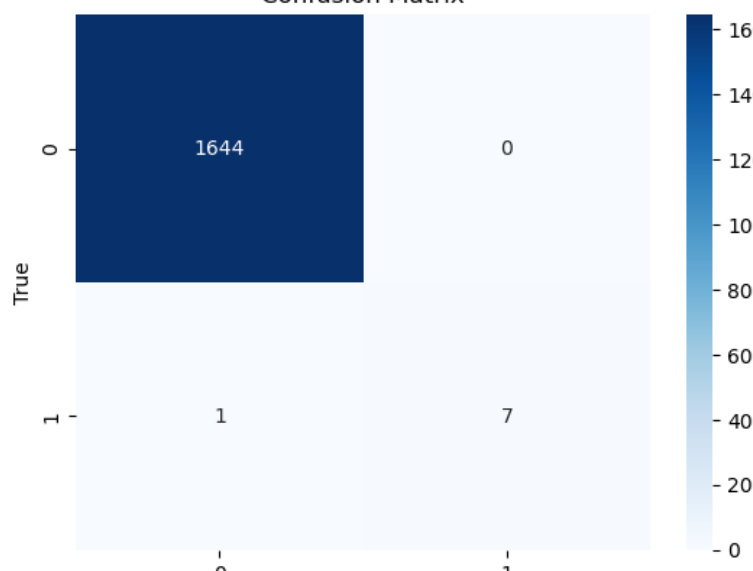
Using randomized search and cross-validation, I attempted to identify the optimal hyperparameters for Random Forest and Gradient Boost. After tuning, best parameters resulted in an F1 Score (weighted) of 0.84 for our Random Forest model and an F1 Score (weighted) of 0.86 for our XG Boost model. There appear to be barely any improvements (if any) to the models after tuning but Gradient Boost appear to be performing the best out of the models tried.

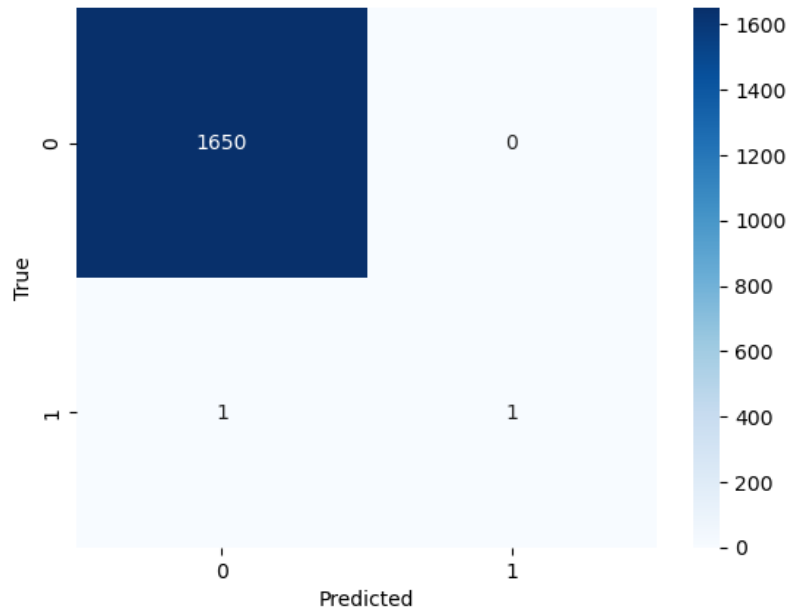
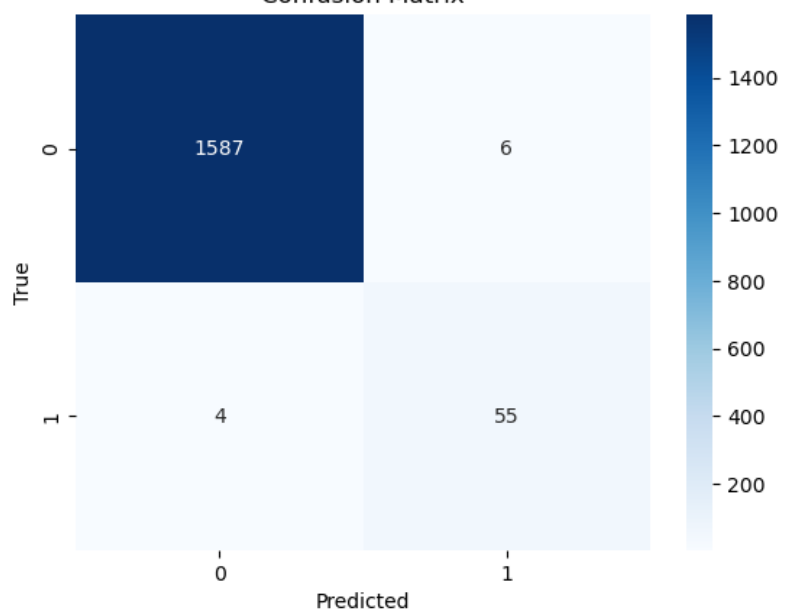
As for models used to predict penalty, I similarly found that ensemble methods largely had superior performance. Models had generally high performance (typically above 0.98). I report the F1 scores (weighted) and confusion matrixes of the best models below for the various penalties of interest:

Table 4: NYPD Penalty Model Performance

NYPD Penalty	Best Model	F1 Score	Confusion Matrix									
Vacation Loss	Random Forest	0.995	<div><p>Confusion Matrix</p><table><tr><th></th><th>0</th><th>1</th></tr><tr><th>True 0</th><td>1598</td><td>1</td></tr><tr><th>True 1</th><td>7</td><td>46</td></tr></table></div>		0	1	True 0	1598	1	True 1	7	46
	0	1										
True 0	1598	1										
True 1	7	46										
Command Discipline A	Gradient Boost	0.986	<div><p>Confusion Matrix</p><table><tr><th></th><th>0</th><th>1</th></tr><tr><th>True 0</th><td>1561</td><td>2</td></tr><tr><th>True 1</th><td>20</td><td>69</td></tr></table></div>		0	1	True 0	1561	2	True 1	20	69
	0	1										
True 0	1561	2										
True 1	20	69										

Command Discipline B	Gradient Boost	0.999	<p>Confusion Matrix</p>  <table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1629</td><td>0</td></tr><tr><th>True 1</th><td>1</td><td>22</td></tr></table>		Predicted 0	Predicted 1	True 0	1629	0	True 1	1	22
	Predicted 0	Predicted 1										
True 0	1629	0										
True 1	1	22										
Formalized Training	Gradient Boost	0.984	<p>Confusion Matrix</p>  <table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1536</td><td>7</td></tr><tr><th>True 1</th><td>18</td><td>91</td></tr></table>		Predicted 0	Predicted 1	True 0	1536	7	True 1	18	91
	Predicted 0	Predicted 1										
True 0	1536	7										
True 1	18	91										

Instructions	Random Forest	0.982	<div>Confusion Matrix</div>  <table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1578</td><td>1</td></tr><tr><th>True 1</th><td>26</td><td>47</td></tr></table>		Predicted 0	Predicted 1	True 0	1578	1	True 1	26	47
	Predicted 0	Predicted 1										
True 0	1578	1										
True 1	26	47										
Left Force	Gradient Boost	0.999	<div>Confusion Matrix</div>  <table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1644</td><td>0</td></tr><tr><th>True 1</th><td>1</td><td>7</td></tr></table>		Predicted 0	Predicted 1	True 0	1644	0	True 1	1	7
	Predicted 0	Predicted 1										
True 0	1644	0										
True 1	1	7										

Other Penalty	Random Forest	0.999	<div><p>Confusion Matrix</p><table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1650</td><td>0</td></tr><tr><th>True 1</th><td>1</td><td>1</td></tr></table></div>		Predicted 0	Predicted 1	True 0	1650	0	True 1	1	1
	Predicted 0	Predicted 1										
True 0	1650	0										
True 1	1	1										
No Penalty	Gradient Boost	0.993	<div><p>Confusion Matrix</p><table><tr><th></th><th>Predicted 0</th><th>Predicted 1</th></tr><tr><th>True 0</th><td>1587</td><td>6</td></tr><tr><th>True 1</th><td>4</td><td>55</td></tr></table></div>		Predicted 0	Predicted 1	True 0	1587	6	True 1	4	55
	Predicted 0	Predicted 1										
True 0	1587	6										
True 1	4	55										

Limitations and Challenges

There are some limitations given preprocessing choices and model selection. The biggest limitation is the imbalance in the dataset. While I have attempted to address it using SMOTE, using SMOTE can increase the risk of overfitting. Simultaneously, given that SMOTE is generating synthetic new data points based on existing data, it might also potentially

increase bias. In the dataset, the target variable 'CCRB Complaint Disposition' had a relatively even distribution of samples for 'Substantiated' and 'Unsubstantiated' classes but lower numbers for 'Exonerated' and 'Unfounded.' Based on confusion matrix results, it appears that SMOTE was not able to fully address the imbalance. Models were generally better at predicting the majority classes and performed poorly on minority classes. Dichotomizing this variable may lead to better overall performance in the future.

As for NYPD Officer Penalty, SMOTE appears to have been more successful. Model performance was high for all penalty types, though confusion matrix results suggest that the models may be more prone to false negatives and less prone to false positives. Models for NYPD Officer Penalty thus may have low recall but high precision. However, the preprocessing approach may have introduced error. In this dataset, officers may receive multiple penalties or only singular penalties. For preprocessing, because we are building models to predict the likelihood of each penalty independently, it doesn't take into account other penalties that are also applying. Here, my methodology assume that each penalty is independent to the others and the models thus don't account for co-occurring penalties. This can be problematic if penalties are linked together and the models also aren't able to predict the actual penalties applied, only the likelihood that a penalty type may be applied for a case. Classifier Chains may be able to address this issue but implementing Multilabel SMOTE has proven to be a challenge to preprocessing.

Other challenges include the computational cost. In this case, ensemble method models had the highest performance, which wasn't unexpected. Hyperparameter tuning was too costly to perform via grid search so I had to opt for Random Search. Other limitations include the issue of (1) temporality and (2) complaint clustering. Current, these models don't account for the time of each complaint and how past complaints can affect future complaints. Simultaneously, it is unclear how the clustering of multiple allegations under singular complaints is affecting results.

References

- About NYPD*. (n.d.). Retrieved January 30, 2025, from <https://www.nyc.gov/site/nypd/about/about-nypd/about-nypd-landing.page>
- Chermak, S., McGarrell, E., & Gruenewald, J. (2006). Media coverage of police misconduct and attitudes toward police. *Policing: An International Journal of Police Strategies & Management*, 29(2), 261–281. <https://doi.org/10.1108/13639510610667664>
- Civilian Complaint Review Board: Allegations Against Police Officers | NYC Open Data*. (n.d.). Retrieved January 31, 2025, from https://data.cityofnewyork.us/Public-Safety/Civilian-Complaint-Review-Board-Allegations-Against/6xgr-kwjq/about_data
- Civilian Complaint Review Board (CCRB) Database*. (n.d.). Retrieved January 31, 2025, from https://data.cityofnewyork.us/browse?Data-Collection_Data-Collection=CCRB%20Complaints%20Database
- Civilian Complaint Review Board: Complaints Against Police Officers | NYC Open Data*. (n.d.). Retrieved January 31, 2025, from https://data.cityofnewyork.us/Public-Safety/Civilian-Complaint-Review-Board-Complaints-Against/2mby-ccnw/about_data
- Civilian Complaint Review Board: Penalties | NYC Open Data*. (n.d.). Retrieved January 31, 2025, from https://data.cityofnewyork.us/Public-Safety/Civilian-Complaint-Review-Board-Penalties/keep-pkmh/about_data
- Crime Stats—Historical—NYPD*. (n.d.). Retrieved March 2, 2025, from <https://www.nyc.gov/site/nypd/stats/crime-statistics/historical.page>
- Cubitt, T. I. C., Gaub, J. E., & Holtfreter, K. (2022). Gender differences in serious police misconduct: A machine-learning analysis of the New York Police Department (NYPD). *Journal of Criminal Justice*, 82, 101976. <https://doi.org/10.1016/j.jcrimjus.2022.101976>
- Dowler, K., & Zawilski, V. (2007). Public perceptions of police misconduct and discrimination: Examining the impact of media consumption. *Journal of Criminal Justice*, 35(2), 193–203. <https://doi.org/10.1016/j.jcrimjus.2007.01.006>
- Harris, C. J., & Worden, R. E. (2014). The Effect of Sanctions on Police Misconduct. *Crime & Delinquency*, 60(8), 1258–1288. <https://doi.org/10.1177/0011128712466933>
- Headley, A. M., D'Alessio, S. J., & Stolzenberg, L. (2020). The Effect of a Complainant's Race and Ethnicity on Dispositional Outcome in Police Misconduct Cases in Chicago. *Race and Justice*, 10(1), 43–61. <https://doi.org/10.1177/2153368717726829>
- Mapping Police Violence*. (n.d.-a). Mapping Police Violence. Retrieved January 31, 2025, from <https://mappingpoliceviolence.org/methodology>
- Mapping Police Violence*. (n.d.-b). Airtable. Retrieved January 31, 2025, from <https://airtable.com/appzVzSeINK1S3EVR/shroOenW19l1m3w0H/tblxearKzw8W7ViN8>
- Rozema, K., & Schanzenbach, M. (2019). Good Cop, Bad Cop: Using Civilian Allegations to Predict Police Misconduct. *American Economic Journal: Economic Policy*, 11(2), 225–268.
- Tapp, S. N., & Davis, E. J. (2020). *Contacts Between Police and the Public, 2020*. US Department of Justice Office of Justice Programs Bureau of Justice Statistics.

Wright II, J. E. (2020). Will They Even Hear Me? How Race Influences Citizen Complaint Outcomes. *Public Performance & Management Review*, 43(2), 257–277.
<https://doi.org/10.1080/15309576.2019.1660188>